

# **Supervised Learning : Linear Regression and Logistic Regression**

Dr. Srijith P K


Computer Science and Engineering

IIT Hyderabad




आई आई टी हैदराबाद  
IIT Hyderabad

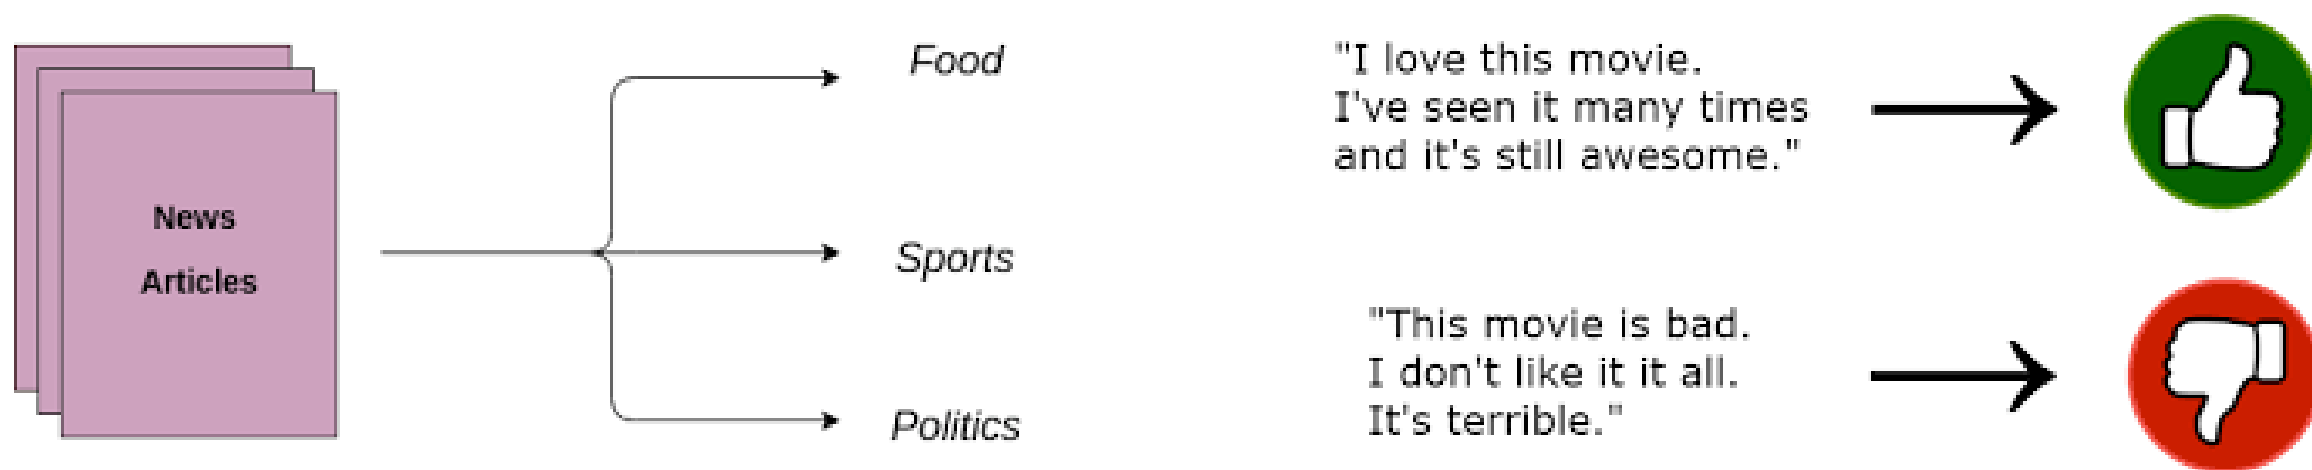

# Supervised learning



# Supervised learning



# Supervised Learning

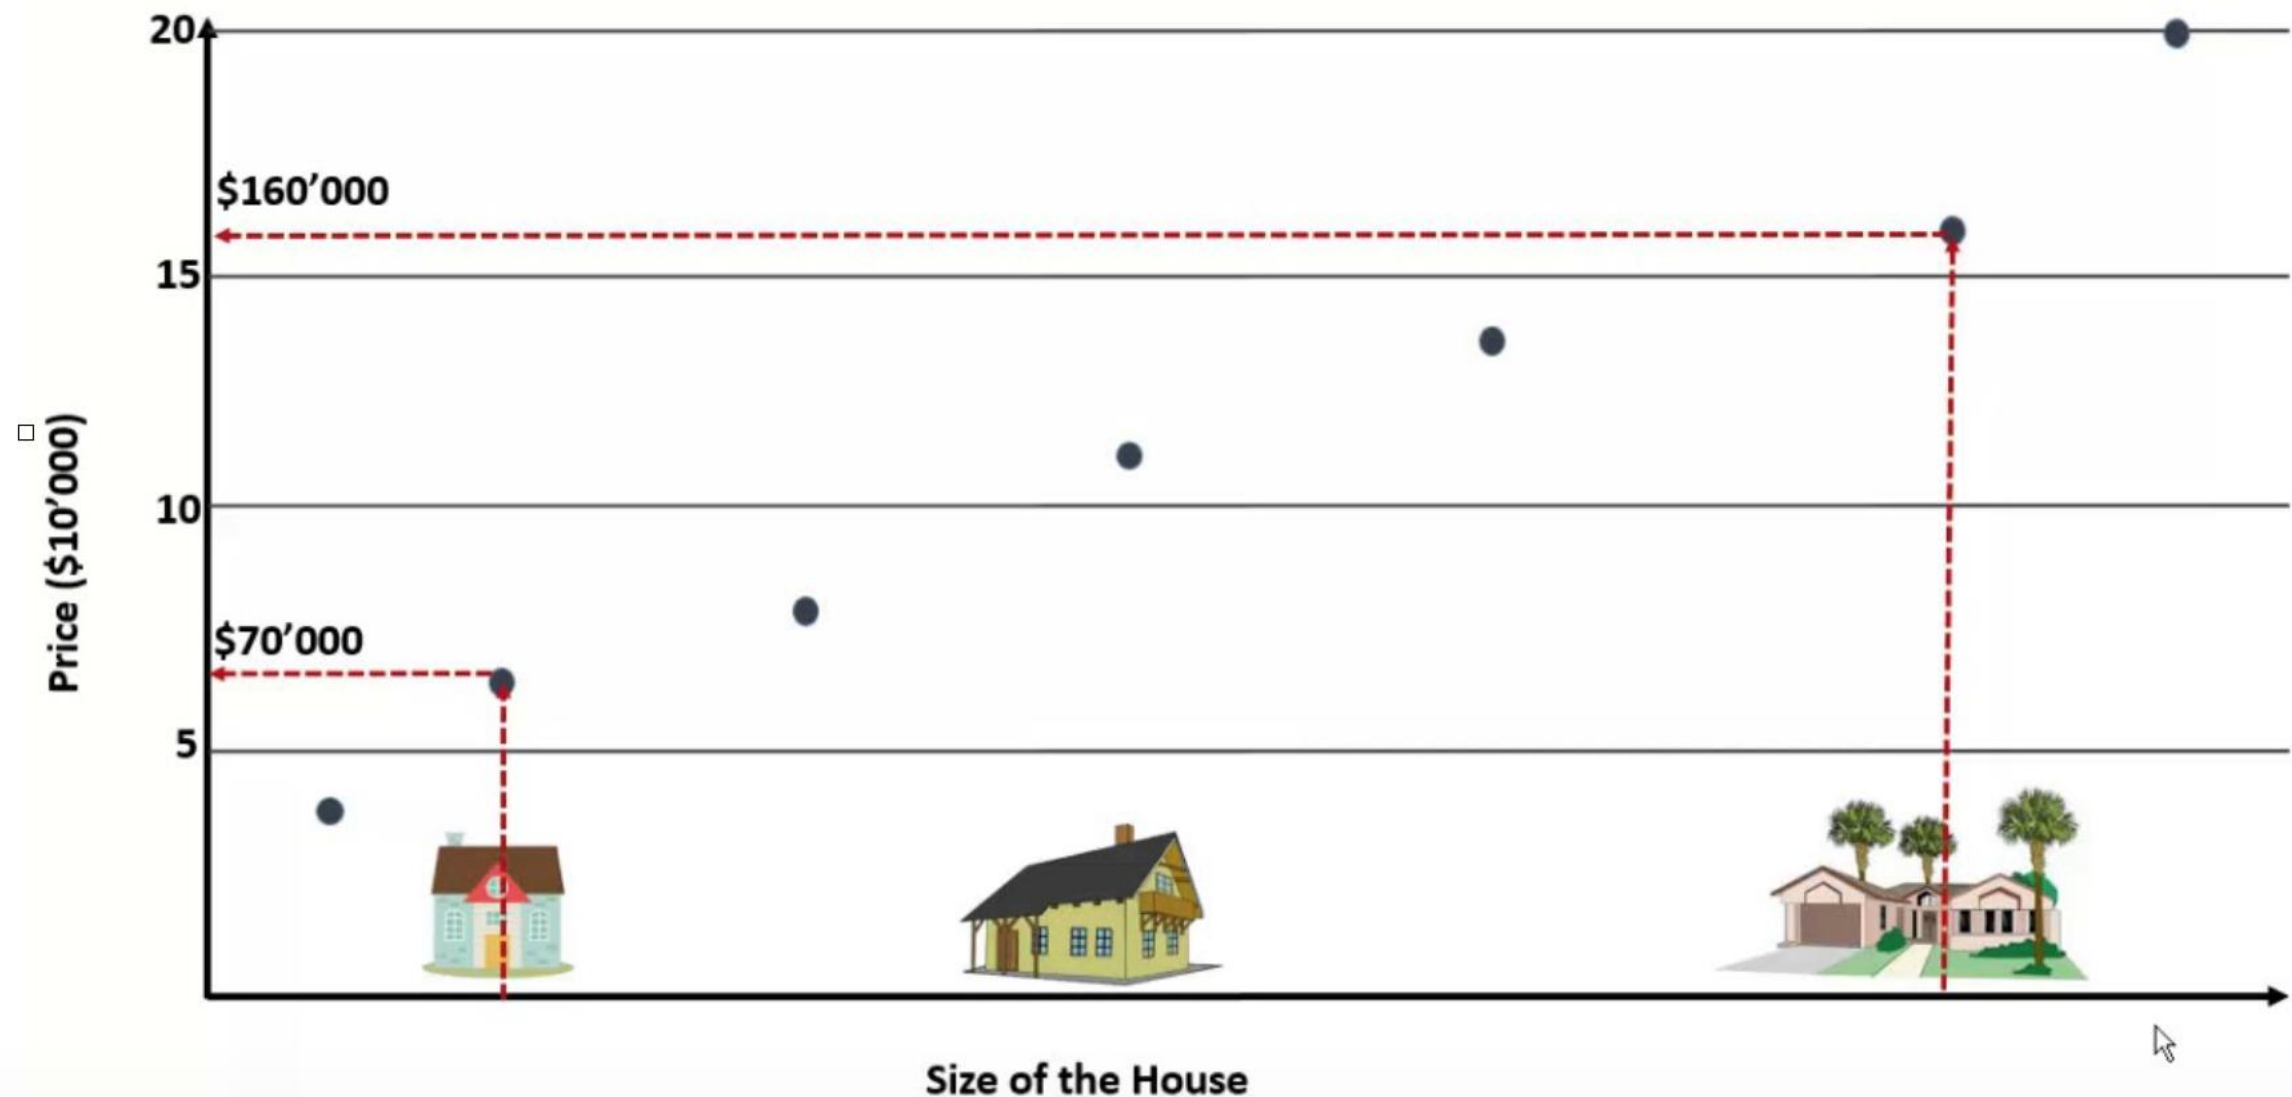


# Outline


- Supervised Learning
  - Regression
    - Linear Regression
    - Logistic Regression
    - Poisson Regression
  - Classification

# Supervised learning : Regression


## Estimating Price of a house



# Supervised learning : Regression



# Supervised learning : Regression



Real valued targets (outputs)

**Generalization performance**

Goal is to learn a function which maps inputs to outputs so that it will **predict well on future data points**




# Airquality data.

- Data set has various air quality parameters in New York city.
- These are the parameters in the data set:
- Daily temperature from May to August
- Solar radiation data
- Ozone data
- Wind data
- Goal : predict the temperature for a particular month in New York using solar radiation, ozone and wind data.


# Airquality data

##		Ozone	<u>Solar.R</u>	Wind	Temp	Month	Day
##	1	41	<u>190</u>	<u>7.4</u>	67	5	1
##	2	36	118	8.0	72	5	2
##	3	12	149	12.6	74	5	3
##	4	18	313	11.5	62	5	4
##	5	NA	<u>NA</u>	14.3	56	5	5
##	6	28	NA	14.9	66	5	6
##	7	23	299	8.6	65	5	7
##	8	19	99	13.8	59	5	8
##	9	8	19	20.1	61	5	9
##	10	NA	194	8.6	69	5	10

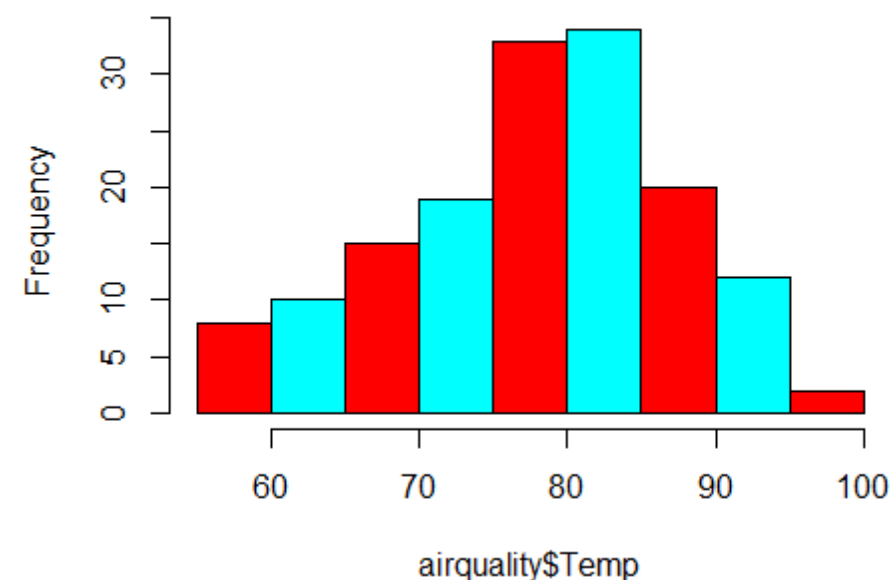
Month 5





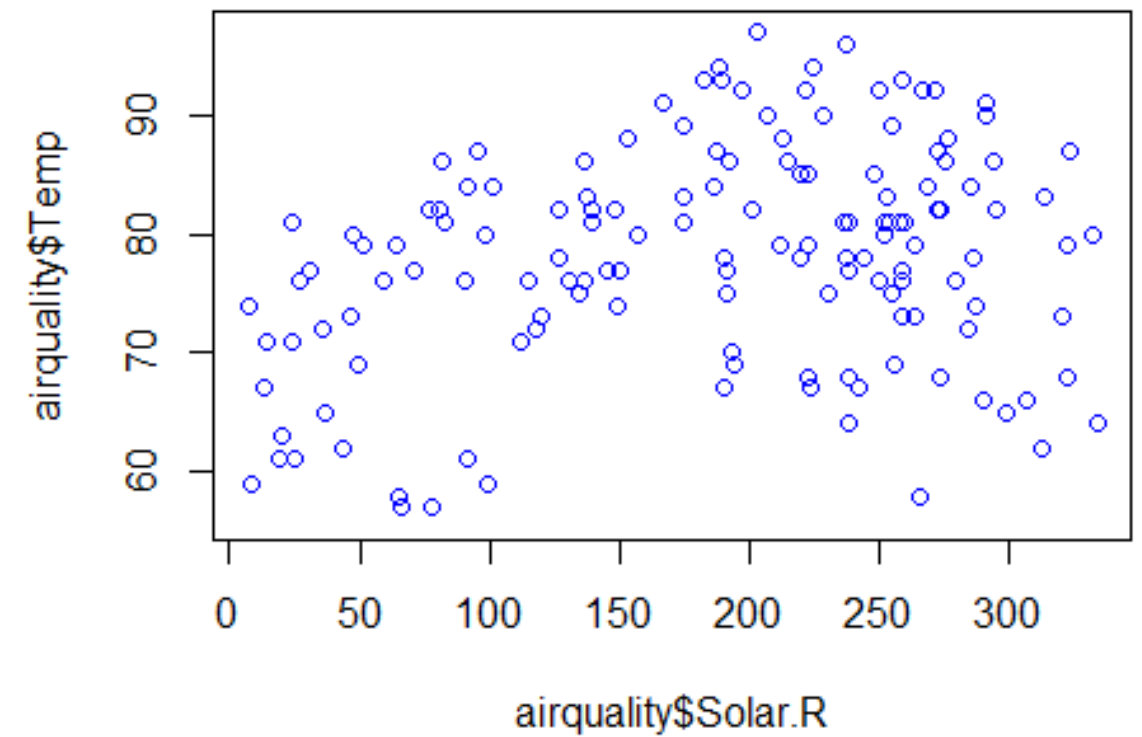

Month 6



Histogram of airquality\$Temp



# Airquality data




# Linear regression

- $\text{Temp} = w_1 \cdot \text{Solar.R} + w_2 \cdot \text{Ozone} + w_3 \cdot \text{Wind} + \text{error}.$
- Temperature of house depends on ozone, wind and solar radiations
- linear regression helps to discover relation between dependent and independent variables

# Linear Regression

- Observations need not lie on a line
- Observations are not generated by a linear line
- Observations are noisy, due to measurement errors



# Linear regression

- Learn a function which maps input to output  $f : X \rightarrow Y$
- Consider a Linear function

Regression Output is real and scalar,  $y \in \mathbb{R}$


Estimated (or predicted) Y value for observation i

Estimate of the regression intercept

Estimate of the regression slope

Value of X for observation i

$$\hat{Y}_i = w_0 + w_1 X_i$$



# Linear regression

- Learn a function which maps input to output  $f: X \rightarrow Y$

Regression Output is real and scalar,  $y \in \mathbb{R}$

- Consider a Linear function


Estimated (or predicted) Y value for observation i

Estimate of the regression intercept

Estimate of the regression slope

Value of X for observation i

$$\hat{Y}_i = w_0 + w_1 X_i$$
$$= X_i^T w$$



1 dim input  $X_i = [1, X_i]^T$

D dim input  $X_i = [1, X_{i1}, \dots, X_{iD}]^T$

$w = [w_0, w_1]^T$

$w = [w_0, w_1, \dots, w_D]^T$


# Linear Regression - Learning Parameters

- Learn the function which passes through as many points as possible :  
Minimize the **Least Squares Error**

$$X_i = [1, X_{i1}, \dots, X_{iD}]^T$$

Design matrix  $X = \begin{pmatrix} X_1 & X_2 & \dots & X_N \end{pmatrix}$   $(D+1) \times N$

$$\begin{aligned} E(w) &= \frac{1}{2} \sum_{i=1}^N (y_i - \hat{y}_i)^2 \\ &= \frac{1}{2} \sum_{i=1}^N (y_i - X_i^T w)^2 \\ &= \frac{1}{2} \| (y - X^T w) \|^2 \end{aligned}$$





# Linear Regression - Learning Parameters


- Learn the function which passes through as many points as possible :  
Minimize the **Least Squares Error**

$$\begin{aligned} E(w) &= \frac{1}{2} \sum_{i=1}^N (y_i - \hat{y}_i)^2 \\ &= \frac{1}{2} \sum_{i=1}^N (y_i - X_i^\top w)^2 \\ &= \frac{1}{2} \| (y - X^\top w) \|^2 \end{aligned}$$

$$\nabla E(w) = Xy - XX^\top w = 0$$

$$X_i = [1, X_{i1}, \dots, X_{iD}]^\top$$

Design matrix  $X = \begin{pmatrix} X_1 & X_2 & \dots & X_N \end{pmatrix}$   $(D+1) \times N$



$$w_{ML} = (XX^\top)^{-1}Xy$$




# Linear Regression - Learning Parameters

- Learn the function which passes through as many points as possible  
: Minimize the **least squares error** using **gradient descent**

$$\nabla E(w) = Xy - XX^T w = 0$$

$$Error_{(m,c)} = \frac{1}{N} \sum_{i=1}^N (y_i - (mx_i + c))^2$$

$$w^{(\tau+1)} = w^{(\tau)} - \eta \nabla E_n$$



# Question


- write down three equations for the line  $y = mx + c$  to go through  $y = 7$  at  $x = -1$ ,  $y = 7$  at  $x = 1$  and  $y = 21$  at  $x = 2$ . Find the least squares solution  $(c, m)$ ?
- Implement In python least squares solution to linear regression
  - Analytical approach
  - Gradient descent approach

## Question

- write down three equations for the line  $y = mx + c$  to go through  $y = 7$  at  $x = -1$ ,  $y = 7$  at  $x = 1$  and  $y = 21$  at  $x = 2$ . Find the least squares solution  $(c, m)$ ?
- Answer : (9,4)

# Non Linear Regression - curve fitting


- Remember high school maths !
- Real-valued target variable  $t$ .
- Training set comprising  $N$  observations



# Regression - curve fitting

- $M$  is the order of the polynomial,  $y(x, w)$  is a nonlinear function of  $x$ , it is a linear function of the coefficients  $w$ .
- Functions, such as the polynomial, which are linear in the unknown parameters have important properties and are called **linear models**

$$y(x, w) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j$$




# Regression - curve fitting

$$y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j$$





- Coefficients will be determined by fitting the polynomial to the training data. This can be done by minimizing an error function

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2$$



# Regression - curve fitting

- Model selection (choosing  $M$ ) : higher order polynomial ( $M = 9$ ), provide excellent fit to the training data but gives a very poor representation of the function









# Regression - curve fitting

- Model selection (choosing  $M$ ) : high degree polynomial ( $M = 9$ ), provide excellent fit to the training data, but gives a very poor representation of the function

**Overfitting**





# Regression - curve fitting

- Generalization performance : root mean square error on test data
- Weights coefficients for  $M=9$  is extremely large !

$$E_{\text{RMS}} = \sqrt{2E(\mathbf{w}^*)/N}$$



$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2$$

	$M = 0$	$M = 1$	$M = 6$	$M = 9$
$w_0^*$	0.19	0.82	0.31	0.35
$w_1^*$		-1.27	7.99	232.37
$w_2^*$			-25.43	-5321.83
$w_3^*$			17.37	48568.31
$w_4^*$				-231639.30
$w_5^*$				640042.26
$w_6^*$				-1061800.52
$w_7^*$				1042400.18
$w_8^*$				-557682.99
$w_9^*$				125201.43



# Regression - curve fitting

- Given model complexity, the over-fitting problem become less severe as the size of the data set increases.



# Curve fitting - regularization



- Add a penalty term to the error function (1.2) in order to discourage the coefficients from reaching large values

$$\tilde{E}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2$$




- Ridge regression : L2 norm

$$\|\mathbf{w}\|^2 \equiv \mathbf{w}^T \mathbf{w} = w_0^2 + w_1^2 + \dots + w_M^2$$

**Regularization  
constant**



# Curve fitting - regularization



	$\ln \lambda = -\infty$	$\ln \lambda = -18$	$\ln \lambda = 0$
$w_0^*$	0.35	0.35	0.13
$w_1^*$	232.37	4.74	-0.05
$w_2^*$	-5321.83	-0.77	-0.06
$w_3^*$	48568.31	-31.97	-0.05
$w_4^*$	-231639.30	-3.89	-0.03
$w_5^*$	640042.26	55.28	-0.02
$w_6^*$	-1061800.52	41.32	-0.01
$w_7^*$	1042400.18	-45.95	-0.00
$w_8^*$	-557682.99	-91.53	0.00
$w_9^*$	125201.43	72.68	0.01

# Regularized Least Squares

- Parameter shrinkage, weight decay

- Ridge regression**  $q=2$  
$$\frac{\lambda}{2} \sum_{j=1}^M w_j^2$$








- Lasso regression**  $q=1$ , if  $\lambda$  is sufficiently large, some of the coefficients are driven to zero

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{t_n - \mathbf{w}^T \phi(\mathbf{x}_n)\}^2 + \frac{\lambda}{2} \sum_{j=1}^M |w_j|$$

- Elastic net regularization

$$\frac{1}{n} \|Y - X\mathbf{w}\|_2^2 + \lambda_1 \sum_{j=1}^d |w_j| + \lambda_2 \sum_{j=1}^d |w_j|^2$$

$$\frac{1}{2} \sum_{n=1}^N \{t_n - \mathbf{w}^T \phi(\mathbf{x}_n)\}^2 + \frac{\lambda}{2} \sum_{j=1}^M |w_j|^q$$



# Ridge Regression

- Regularized Least Squares

$$\frac{1}{2} \sum_{n=1}^N \{t_n - \mathbf{w}^T \phi(\mathbf{x}_n)\}^2 + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}.$$

$$\Phi = \begin{pmatrix} \phi_0(\mathbf{x}_1) & \phi_1(\mathbf{x}_1) & \cdots & \phi_{M-1}(\mathbf{x}_1) \\ \phi_0(\mathbf{x}_2) & \phi_1(\mathbf{x}_2) & \cdots & \phi_{M-1}(\mathbf{x}_2) \\ \vdots & \vdots & \ddots & \vdots \\ \phi_0(\mathbf{x}_N) & \phi_1(\mathbf{x}_N) & \cdots & \phi_{M-1}(\mathbf{x}_N) \end{pmatrix} \quad N \times M$$

- Show that the regularized least squares solution is

$$\mathbf{w} = (\lambda \mathbf{I} + \Phi^T \Phi)^{-1} \Phi^T \mathbf{t}.$$

Stable and  
unique solution

# Regularized Least Squares : Cross Validation

How to choose  $\lambda$  ?  
Use validation data!

$$\frac{1}{2} \sum_{n=1}^N \{t_n - \mathbf{w}^T \phi(\mathbf{x}_n)\}^2 + \frac{\lambda}{2} \sum_{j=1}^M |w_j|^q$$



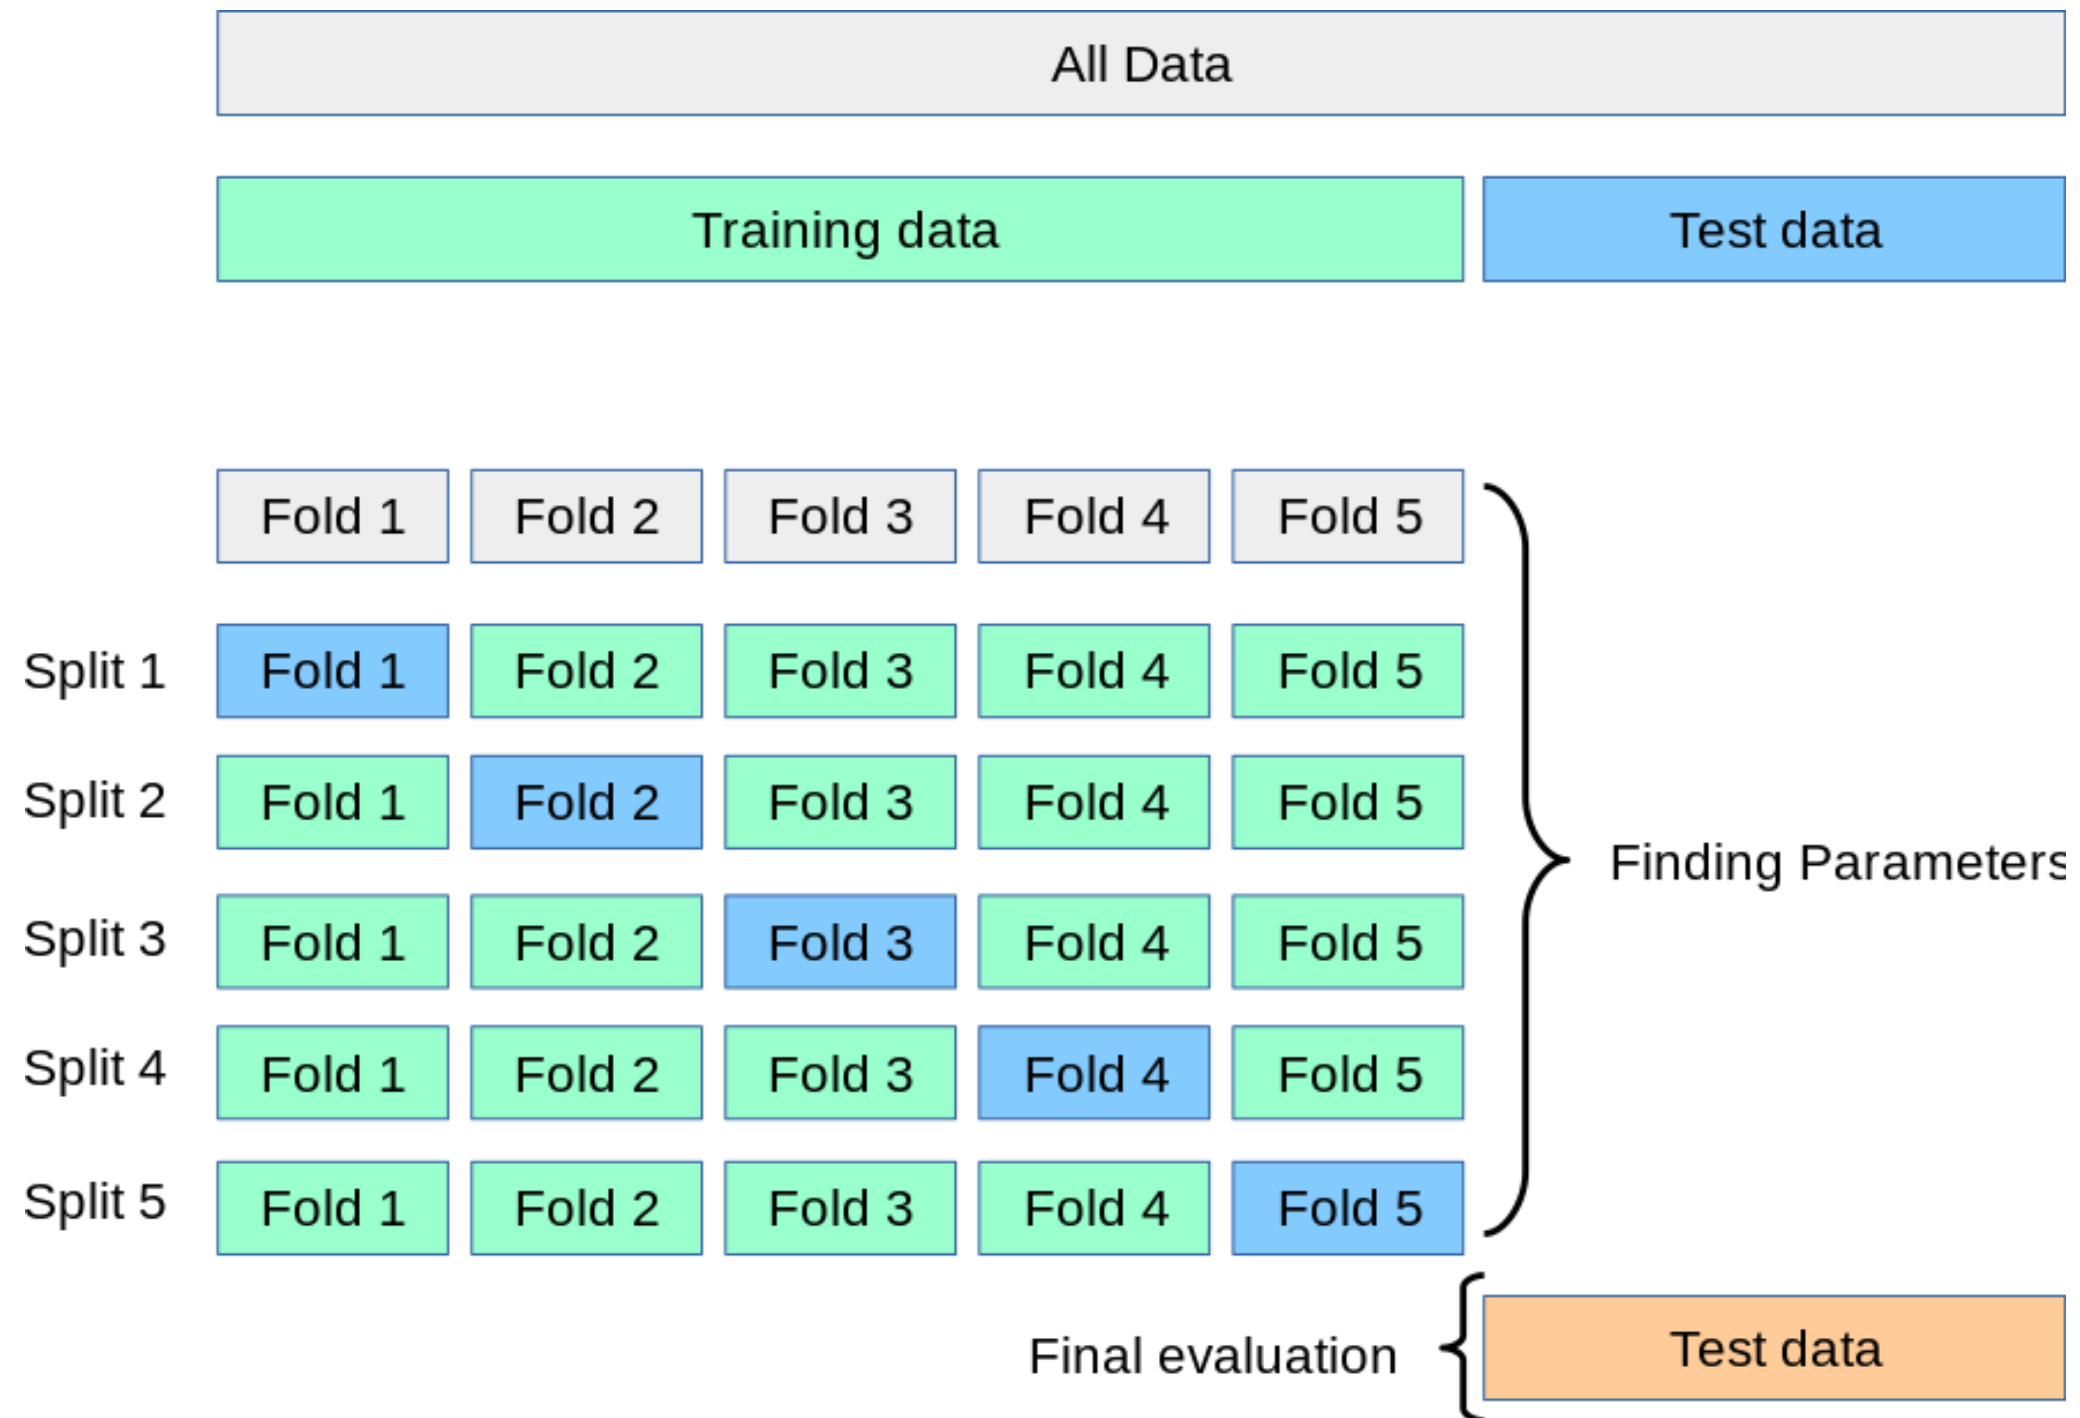
1. **Training set** is a set of examples used for learning a model parameters (e.g., weight vector  $\mathbf{w}$  in linear regression)
2. **Validation set** is a set of examples that cannot be used for learning the model parameter but can help tune model hyper-parameters e.g. Regularization constant in LR. Validation helps control overfitting.
3. **Test set** is used to assess the performance of the final model and provide an estimation of the test error.

Example: Split the data randomly into 60% for training, 20% for validation and 20% for testing.

**Note: Dont use the test set to further tune the parameters or revise the model.**




# Cross Validation



# Probabilistic Interpretation : Least Squares = Maximum likelihood estimation

$$y_i \sim w \cdot x_i + N(0, \sigma^2) = N(w \cdot x_i, \sigma^2)$$




$$L = \prod_i \exp -\frac{1}{2\sigma^2} (X_i^T w - y_i)^2 = \exp -\frac{1}{2\sigma^2} \sum_i (X_i^T w - y_i)^2$$

$$\underset{w}{\operatorname{argmax}} L = \underset{w}{\operatorname{argmin}} E$$


- Similarly Regularized least squares is same as maximum a posteriori estimate assuming  $p(w)$  to be a Gaussian :  $\underset{w}{\operatorname{argmax}} p(y_i | w, x_i) p(w)$

# Supervised Learning : Classification

- Binary classification :  $y = \{0,1\}$
- Multiclass classification :  $y = \{1,2,...K\}$

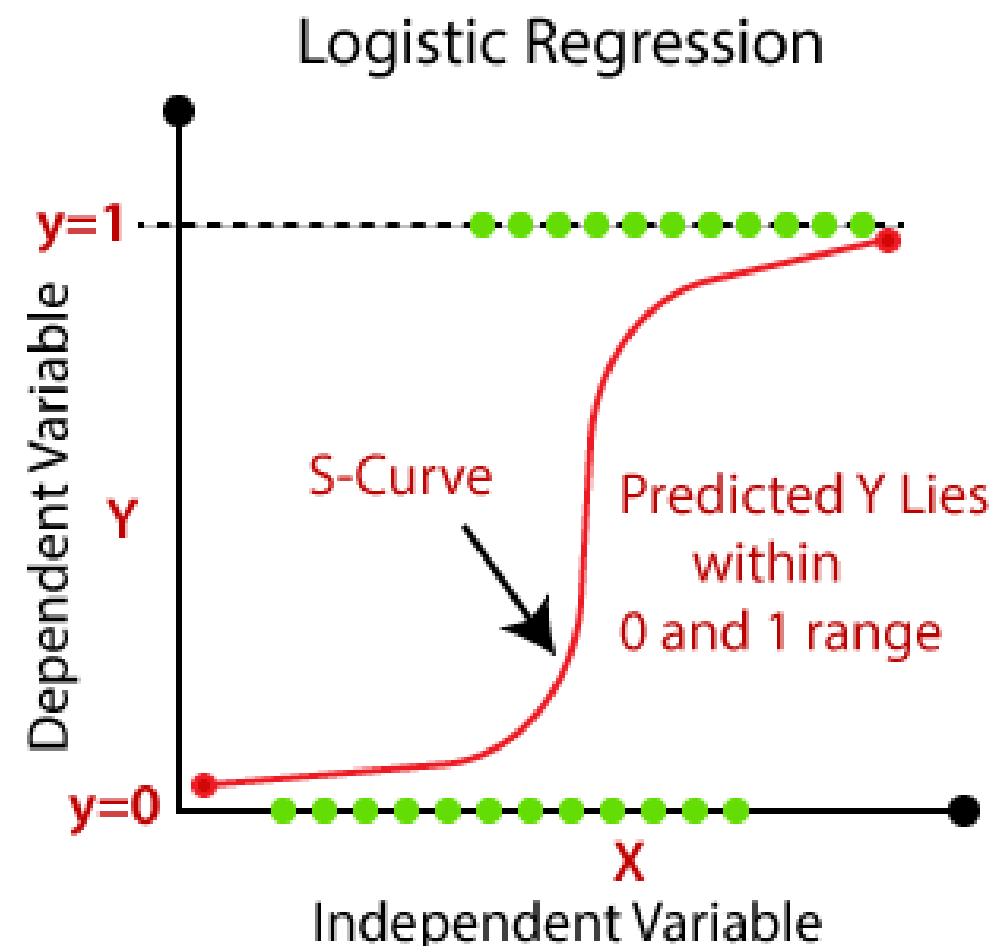



$p(y|x)?$



# Linear regression to Logistic regression



- Y takes value 0 or 1



# Logistic Regression

- A discriminative approach which directly models  $p(y|x)$
- To predict an outcome variable that is categorical from one or more categorical or continuous predictor variables.
- Let  $X$  be the data instance, and  $Y$  be the class label  $\{0,1\}$  :  
Model  $P(Y|X)$  directly using a **Sigmoid function**:

**Logistic Sigmoid :**  $P(Y = 1 | \mathbf{X}) = \frac{1}{1 + e^{-\mathbf{w}\mathbf{x}}}$




# Logistic Regression


- To predict an outcome variable that is categorical from one or more categorical or continuous predictor variables.
- Let  $X$  be the data instance, and  $Y$  be the class label:  
Model  $P(Y|X)$  directly using a **Sigmoid function**:

$$P(Y = 1 | \mathbf{X}) = \frac{1}{1 + e^{-w\mathbf{x}}}$$


**[HW] Find derivative of  $s(w) = p(y=1|X)$ !**



# Logistic Regression



$$P(Y = 1 | \mathbf{X}) = \frac{1}{1 + e^{-\mathbf{w}\mathbf{x}}}$$



# Logistic Regression

- In logistic regression, we learn the conditional distribution  $P(y|x)$
- Let  $p_y(x;w)$  be our estimate of  $P(y|x)$ , where  $w$  is a vector of adjustable parameters.
- Assume there are two classes,  $y = 0$  and  $y = 1$  and

$$p_1(\mathbf{x}; \mathbf{w}) = \frac{1}{1 + e^{-\mathbf{w}\mathbf{x}}} \quad p_0(\mathbf{x}; \mathbf{w}) = 1 - \frac{1}{1 + e^{-\mathbf{w}\mathbf{x}}} = \frac{1}{1 + e^{\mathbf{w}\mathbf{x}}}$$

- This is equivalent to

$$\log \frac{p_1(\mathbf{x}; \mathbf{w})}{p_0(\mathbf{x}; \mathbf{w})} = \mathbf{w}\mathbf{x}$$

- That is, the log odds of class 1 is a linear function of  $x$
- Q: How to find **W**?

- Alternate representation of  $p(y|x)$  :  $p_y(\mathbf{x}; \mathbf{w}) = \frac{1}{1 + e^{-y\mathbf{w}\mathbf{x}}} ; y = \{-1, 1\}$



# Logistic Regression

- Conditional data likelihood - Probability of observed  $Y$  values in the training data, conditioned on corresponding  $X$  values.
- We choose parameters  $w$  that satisfy

$$\mathbf{w} = \arg \max_{\mathbf{w}} \prod_l P(y^l | \mathbf{x}^l, \mathbf{w})$$

- where
  - $\mathbf{w} = \langle w_0, w_1, \dots, w_n \rangle$  is the vector of parameters to be estimated,
  - $y^l$  denotes the observed value of  $Y$  in the  $l$ th training example, and
  - $\mathbf{x}^l$  denotes the observed value of  $\mathbf{X}$  in the  $l$ th training example

# Logistic Regression

- Equivalently, we can work with log of conditional likelihood:

$$\mathbf{w} = \arg \max_{\mathbf{w}} \sum_l \ln P(y^l | \mathbf{x}^l, \mathbf{w})$$

- Conditional data log likelihood,  $l(\mathbf{w})$ , can be written as

$$l(\mathbf{w}) = \sum_l y^l \ln P(y^l = 1 | \mathbf{x}^l, \mathbf{w}) + (1 - y^l) \ln P(y^l = 0 | \mathbf{x}^l, \mathbf{w})$$

- Note here that  $Y$  can take only values 0 or 1, so only one of the two terms in the expression will be non-zero for any given  $y^l$

# Logistic Regression

- We need to estimate:

$$\mathbf{w} = \arg \max_{\mathbf{w}} \sum_l \ln P(y^l | \mathbf{x}^l, \mathbf{w})$$

$$l(\mathbf{w}) = \sum_l y^l \ln P(y^l = 1 | \mathbf{x}^l, \mathbf{w}) + (1 - y^l) \ln P(y^l = 0 | \mathbf{x}^l, \mathbf{w})$$

- Equivalently, we can minimize negative log likelihood using gradient descent technique
- No closed-form solution though. Iterative method required.
- [HW] Find the derivative of  $l(\mathbf{w})$  !

# Logistic Regression

- Overfitting can arise especially when data has very high dimensions and is sparse.
- One approach -> modified “penalized log likelihood function,” which penalizes large values of  $\mathbf{w}$ , as before.

$$\mathbf{w} = \arg \max_{\mathbf{w}} \sum_l \ln P(y^l | \mathbf{x}^l, \mathbf{w}) - \frac{\lambda}{2} \|\mathbf{w}\|^2$$

- [HW] Find the Derivative !

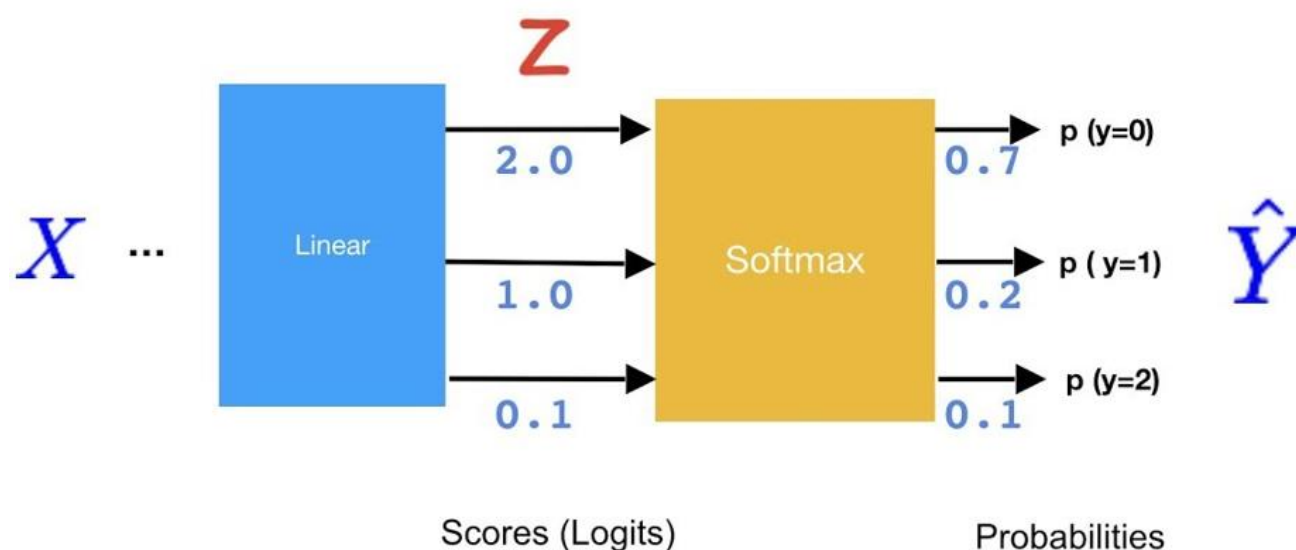
# Logistic Regression

- LR: Functional form of  $P(Y|X)$ , no assumption on  $P(X|Y)$
- LR is a linear classifier
- LR optimized by conditional likelihood
- Extending logistic regression to multiple classes
  - Use softmax for each class  $k$ !

$$p(y = k|x) = \frac{\exp(\mathbf{w}_k^T \mathbf{x})}{\sum_{i=1}^K \exp(\mathbf{w}_i^T \mathbf{x})}$$

Meet Softmax

$$\sigma(\mathbf{z})_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}} \quad \text{for } j = 1, \dots, K.$$




# Classification : Evaluation metrics

**Accuracy :**  $\frac{1}{N} \sum_{i=1}^N \mathbb{I}[y_i == \hat{y}_i]$

		Actual Label	
		Positive	Negative
Predicted Label	Positive	<b>True Positive (TP)</b>	<b>False Positive (FP)</b>
	Negative	<b>False Negative (FN)</b>	<b>True Negative (TN)</b>

<b>Accuracy</b>	$(TP + TN) / (TP + TN + FP + FN)$	The percentage of predictions that are correct
<b>Precision</b>	$TP / (TP + FP)$	The percentage of positive predictions that are correct
<b>Sensitivity (Recall)</b>	$TP / (TP + FN)$	The percentage of positive cases that were predicted as positive
<b>Specificity</b>	$TN / (TN + FP)$	The percentage of negative cases that were predicted as negative

# Supervised learning : Regression



Number of vehicles passing a junction

Y take values 0,1,2,3,..... but not 2.1, 3.4, 5.55.....


$p(y|x)?$

# Poisson Regression

- Poisson distribution : Model number of events occurring in a fixed interval of time/space

$$P(k \text{ events in interval}) = e^{-\lambda} \frac{\lambda^k}{k!}$$

- $\lambda$  is the average (mean) number of events
- $Y$  has a Poisson distribution, and assumes the logarithm of its expected value can be modeled by a linear combination of unknown parameters.



$$\lambda := E(Y \mid x) = e^{\theta' x}$$

$$p(y \mid x; \theta) = \frac{\lambda^y}{y!} e^{-\lambda} = \frac{e^{y\theta' x} e^{-e^{\theta' x}}}{y!}$$



# Poisson Regression : Learning parameters

- Likelihood 
$$p(y_1, \dots, y_m \mid x_1, \dots, x_m; \theta) = \prod_{i=1}^m \frac{e^{y_i \theta' x_i} e^{-e^{\theta' x_i}}}{y_i!}.$$
- Estimate parameters by maximum likelihood estimation

$$\ell(\theta \mid X, Y) = \log L(\theta \mid X, Y) = \sum_{i=1}^m \left( y_i \theta' x_i - e^{\theta' x_i} - \log(y_i!) \right).$$

- Use gradient descent to find the optimal value of  $\theta$ .

# Thank you !

## Reference

- [1] Christopher Bishop, Pattern Recognition and Machine Learning**
- [2] Kevin Murphy, Machine Learning : A Probabilistic Perspective**